

pyAmpliCol 0.1.0

Methodology and Performance Report

Rikkert Frederix

Valentin Hirschi

Timea Vitos

Release 0.1.0

Abstract

This report defines the public theory, validation, and performance methodology for `pyAmpliCol`, release 0.1.0. It fixes the canonical process coverage for the built-in Standard Model, the same physics loaded as a UFO/JSON Standard Model (UFO-SM), a scalar contact model, and a scalar-gravity model. Its tables present reproducible measurements tied to an effective configuration, execution environment, generated artifact, and independent-reference provenance.

Contents

1	Scope	2
2	Current Recursion and Model Lowering	2
3	Colour and Runtime Semantics	3
4	Benchmark Methodology	3
5	Standard-Model Process Matrices	5
6	Z-plus-Jets Evaluator Ladders	20
7	Colour-Singlet Model Ladders	23
8	Interpretation	23
9	Worked $d\bar{d} \rightarrow Zgg$ Example	24
10	The $d\bar{d} \rightarrow Zgg$ DAG and Rusticol	27
11	Arbitrary UFO and Serialized-JSON Models	31

1 Scope

pyAmpliCol generates tree-level scattering-amplitude artifacts and evaluates them through the Rusticol runtime. The Python layer resolves models and processes, constructs a shared-current directed acyclic graph (DAG), lowers bounded stages to Symbolica evaluators, and writes a versioned process artifact. The runtime loads that artifact and provides summed and resolved evaluation through one physics contract to Python, Rust 2021, C++17, and Fortran 2008 clients.

The report has three purposes:

1. specify the observables and conditions required for a publishable benchmark;
2. preserve a stable matrix of representative process families and colour approximations; and
3. separate measured data from prose by rendering every result table from validated measurement records.

An `N/A` entry means that the current result set has no recorded measurement for that slot. It is neither a zero, a failure, nor a value carried over from another implementation. Explicit failure labels such as `T/O`, `RAM>LIMIT`, `VALIDATION FAILED`, and `ERROR` identify a recorded timeout, resource limit, numerical mismatch, or execution failure rather than a physics value.

2 Current Recursion and Model Lowering

2.1 Shared-current DAG

For an external subset I , particle state t , and local spin or flow label χ , an off-shell current has the schematic recursion

$$J_{t,\chi}^\alpha(I) = P^\alpha_\beta(t, p_I) \sum_{I=A \dot{\cup} B} V^\beta_{\gamma\delta}(t_A, t_B \rightarrow t) J_{t_A, \chi_A}^\gamma(A) J_{t_B, \chi_B}^\delta(B), \quad p_I = \sum_{i \in I} p_i. \quad (1)$$

Current identity includes the particle, external labels, helicity ancestry, chirality or spin state, flavour and charge flow, colour state, momentum mask, and any auxiliary-current kind. Contributions with the same complete identity share one node. A topological sweep therefore evaluates each retained current once per phase-space point and reuses it in all dependent amplitudes.

External UFO or serialized JSON sources are compiled to the same canonical model representation as the process generator consumes. Particle metadata, interactions, coupling orders, Lorentz tensors, colour tensors, propagators, and runtime parameters come from that representation. UFO directories contain trusted executable Python; serialized JSON is the portable data-only source.

Optimizations on the external-model path are proof-gated rather than selected by particle names or PDG values. Exact normalized colour and Lorentz expressions, propagator metadata, and coupling relations certify current reflection, helicity, and reusable-kernel identities. For overlapping Standard-Model validation processes, the resulting current, interaction, amplitude-root, and colour-sector counts must match the built-in compatibility model after applying the same independently inferred perturbative-order envelope. An external compiler may retain fewer evaluator calls only when an additional exact kernel-equivalence certificate proves that reuse.

2.2 Contact terms and spin two

A vertex with valence above three is lowered to a deterministic tree of non-propagating auxiliary currents. The construction inserts the original coupling once and introduces no pole. Direct contraction of the original contact tensor is the validation oracle for the scalar-contact ladder.

The scalar-gravity ladder exercises rank-two external states and momentum-dependent interactions. For a massless graviton, the helicity-two polarization tensor is represented as

$$\epsilon_{\pm 2}^{\mu\nu} = \epsilon_\pm^\mu \epsilon_\pm^\nu. \quad (2)$$

Validation checks symmetry, transversality, tracelessness, normalization, and agreement between ordinary and higher-precision evaluation.

3 Colour and Runtime Semantics

The process matrices use three colour contracts:

LC	Leading-colour ordered amplitudes with physical colour-flow selectors.
NLC	The leading term plus the first retained $1/N_c^2$ -suppressed contribution.
Full	The complete supported sparse colour metric for the generated amplitude basis.

The optimized `Runtime.evaluate()` path returns one summed matrix element per point.

The resolved method `Runtime.evaluate_resolved()` exposes LC values with axes (`point`, `helicity`, `flow`). NLC and full-colour artifacts expose one already contracted colour component per helicity. Summing all non-point axes must reproduce the optimized total. Public selectors use physical helicity and flow identifiers; graph-local colour-sector identifiers are not report inputs.

All language bindings load the same process manifest and evaluator payloads. Cross-language agreement therefore validates the runtime interfaces and data layout. A separately generated reference is required for an independent physics comparison.

The default JIT artifact embeds a self-contained SymJIT application for f64 evaluation. Rusticol executes that payload through the separate MIT-licensed SymJIT runtime without importing Symbolica or consulting its runtime license state. Compiled ASM and C++ f64 payloads are likewise loaded directly by the native core. Python evaluation above f64 precision uses the retained Symbolica evaluator state and therefore requires Symbolica at runtime.

4 Benchmark Methodology

4.1 Configuration contract

Each measurement begins from an immutable resolved configuration. Unless a table variant says otherwise, the report contract uses the release defaults: requested SymJIT optimization level 3, evaluator batch size 128, output chunk size 128, stage-local parameter layout, ten Horner iterations, backend-default common-pair elimination, a ten-second benchmark target, two warmup runs, and at least five timed samples. The process matrices always request level 3; only explicitly labelled comparison rows, such as JIT level 1 in the dedicated *Z* ladder, override it. SymJIT may apply its documented backend fallback independently of the requested level. Requested and effective configurations are retained separately so resource or license adjustments remain visible.

Generation time is wall time for one explicit `Generator.generate()` operation starting from an already compiled canonical model and ending with a committed, loadable process artifact. It includes process expansion, DAG and evaluator construction, evaluator compilation, validation-point preparation, and API-bundle emission. Loading and compiling the UFO/JSON or built-in model is performed before the generation timer, passed to either ordinary generation or the selected `GenerationSlice`, and recorded separately as `model_precompile_seconds`. It is not included in a process-generation cell. Existing artifacts may be benchmarked, but their generation slot remains N/A unless a matching generation measurement is available.

Runtime measurements use `BenchmarkRunner` and the summed `Runtime.evaluate()` observable. The reported wall time begins inside Rusticol after caller-language array conversion and covers repeated ordinary core evaluations, including source preparation, stage input packing, evaluator calls, output assignment, and reduction. The evaluator time is the subset spent inside the generated stage and amplitude evaluators. Python/NumPy packing and language-binding conversion are therefore excluded from both. A measurement records sample count, time per point, uncertainty, effective settings, and environment. Comparisons are valid only for the same process, colour accuracy, parameters, phase-space ensemble, precision, selectors, and normalization.

4.2 Benchmark platform

Campaign metadata pending. The table-filling campaign must replace this placeholder with the exact benchmark architecture before publication: host or cluster partition, CPU model and microarchitecture, instruction-set features, physical/logical core topology, installed memory, operating system and kernel, Python/Rust/C/C++/Fortran toolchain versions, and the process-worker and evaluator-thread policy used for generation and timing. Measurements in different tables must identify any platform or policy change explicitly.

4.3 Numerical validation

The matrix element and its independent reference are evaluated at the same deterministic points with the same parameter card. Shared-core language checks use relative tolerance 10^{-12} and absolute tolerance 10^{-15} . Independent reference comparisons use relative tolerance 10^{-8} , with higher precision used to diagnose unstable points. A populated result must also satisfy the resolved-sum identity and preserve expected structural zeros.

No timing is accepted as a numerical validation result. Conversely, a valid matrix element does not populate generation or runtime fields without a matching timing run.

Result provenance. Each displayed number is retained with its resolved configuration, software revision, hardware and execution environment, numerical check, and timing statistics. A result enters the report only after its record is internally consistent and its pointwise comparison has the status required by the table.

5 Standard-Model Process Matrices

The canonical matrix spans electroweak vector production, charged-current and leptonic channels, heavy-quark production, pure gluons, multi-boson final states, Higgs-associated final states, and three- and four-quark-line cases. Here n is the number of final-state particles. LC coverage retains the declared grid through $n = 9$, with the four bounded families ending at $n = 8$. NLC and full-colour grids retain $n = 1, \dots, 5$; cells below a family’s minimum physical multiplicity remain explicitly marked `N/A`.

The UFO-SM matrices use the same process definitions, colour contracts, and numerical input parameters as the built-in-SM matrices. Their source is the packaged Standard Model loaded through the generic UFO/JSON compiler. The comparison therefore isolates model loading and generic lowering without changing the requested physics or the original Fortran `AmpliCol` reference.

How to read a process-matrix cell

Every process-matrix comparison uses `pyAmpliCol` JIT at requested optimization level 3. The process matrices do not include ASM or C++ measurements; those backends appear in the dedicated Z ladder. For a time T , a parenthesized xr is the dimensionless ratio

$$r = \frac{T_{\text{pyAmpliCol}}}{T_{\text{AmpliCol}}}. \quad (3)$$

Green, orange, and red denote respectively $r < 1$, $1 \leq r < 2$, and $r \geq 2$; thus green means that `pyAmpliCol` is faster. The compact layout is:

contract	line and reference entry	pyAmpliCol entries
LC, upper	<code>AmpliCol</code> generation time in seconds	selected-flow/helicity-sum generation ratio; all-flows/fixed-helicity generation ratio
LC, lower	<code>AmpliCol</code> selected-flow/all-flows runtimes in microseconds per point, separated by “/”	for each workload, (<code>xwall eval</code>) gives the Rusticol-core wall and evaluator-only ratios against the matching <code>AmpliCol</code> runtime
NLC/full, upper	<code>AmpliCol</code> generation time in seconds	<code>pyAmpliCol</code> JIT O3 generation ratio
NLC/full, lower	<code>AmpliCol</code> fully contracted runtime in microseconds per point	(<code>xwall eval</code>) gives the Rusticol-core wall and evaluator-only ratios against that runtime

The two LC generation ratios correspond to separately generated selected-flow and complete-flow `pyAmpliCol` artifacts. `AmpliCol` generates one library that supports both runtime workloads, so the same `AmpliCol` generation time is the denominator of both ratios. “Selected flow, helicity sum” evaluates one matched physical LC ordering and sums all supported helicities. “All flows, fixed helicity” evaluates every represented physical LC ordering for one supported helicity. NLC and full colour instead report one helicity-summed, fully colour-contracted workload; no per-flow scalar is defined for those contracts.

At the foot of each multiplicity column, every summary cell has two lines. The first contains the `AmpliCol` values as `min|max|mean|median`; generation is in seconds and runtime in microseconds per point. The second contains the paired `pyAmpliCol`/`AmpliCol` multipliers as `min|max|mean|median|ratio-of-sums`, where the last field is $\sum_i T_{\text{pyAmpliCol},i} / \sum_i T_{\text{AmpliCol},i}$ over the same valid process pairs. LC provides separate generation and runtime summary rows for its two workloads. NLC and full colour provide generation, core-wall runtime, and evaluator-only runtime rows; the last two share the same `AmpliCol` reference runtime because `AmpliCol` has no corresponding evaluator-only timing split.

5.1 Built-in SM LC process matrix

ID	base process	n=1	n=2	n=3
1	$d\bar{d} \rightarrow Z + (n-1)g$	N/A	N/A	N/A
2	$u\bar{d} \rightarrow W^+ + (n-1)g$	N/A	N/A	N/A
3	$d\bar{d} \rightarrow e^+e^- + (n-2)g$	N/A	N/A	N/A
4	$u\bar{d} \rightarrow e^+\nu_e + (n-2)g$	N/A	N/A	N/A
5	$d\bar{d} \rightarrow ZZ + (n-2)g$	N/A	N/A	N/A
6	$gg \rightarrow t\bar{t} + (n-2)g$	N/A	N/A	N/A
7	$d\bar{d} \rightarrow t\bar{t} + (n-2)g$	N/A	N/A	N/A
8	$gg \rightarrow gg + (n-2)g$	N/A	N/A	N/A
9	$d\bar{d} \rightarrow ZZZ + (n-3)g$	N/A	N/A	N/A
10	$d\bar{d} \rightarrow e^+e^-ZH + (n-4)g$	N/A	N/A	N/A
11	$d\bar{d} \rightarrow t\bar{t}ZH + (n-4)g$	N/A	N/A	N/A
12	$d\bar{d} \rightarrow e^+e^-e^+e^- + (n-4)g$	N/A	N/A	N/A
13	$d\bar{d} \rightarrow u\bar{u}s\bar{s} + (n-4)g$	N/A	N/A	N/A
14	$d\bar{d} \rightarrow u\bar{u}s\bar{s}c\bar{c} + (n-6)g$	N/A	N/A	N/A
gen O3 one-flow, hel. sum		N/A	N/A	N/A
gen O3 all-flows, one-hel		N/A	N/A	N/A
run O3 one-flow, hel. sum		N/A	N/A	N/A
run O3 all-flows, one-hel		N/A	N/A	N/A

Built-in SM LC process matrix (continued)

ID base process	n=4	n=5	n=6
1 $d\bar{d} \rightarrow Z + (n-1)g$	N/A	N/A	N/A
2 $u\bar{d} \rightarrow W^+ + (n-1)g$	N/A	N/A	N/A
3 $d\bar{d} \rightarrow e^+e^- + (n-2)g$	N/A	N/A	N/A
4 $u\bar{d} \rightarrow e^+\nu_e + (n-2)g$	N/A	N/A	N/A
5 $d\bar{d} \rightarrow ZZ + (n-2)g$	N/A	N/A	N/A
6 $gg \rightarrow t\bar{t} + (n-2)g$	N/A	N/A	N/A
7 $d\bar{d} \rightarrow t\bar{t} + (n-2)g$	N/A	N/A	N/A
8 $gg \rightarrow gg + (n-2)g$	N/A	N/A	N/A
9 $d\bar{d} \rightarrow ZZZ + (n-3)g$	N/A	N/A	N/A
10 $d\bar{d} \rightarrow e^+e^-ZH + (n-4)g$	N/A	N/A	N/A
11 $d\bar{d} \rightarrow t\bar{t}ZH + (n-4)g$	N/A	N/A	N/A
12 $d\bar{d} \rightarrow e^+e^-e^+e^- + (n-4)g$	N/A	N/A	N/A
13 $d\bar{d} \rightarrow u\bar{u}s\bar{s} + (n-4)g$	N/A	N/A	N/A
14 $d\bar{d} \rightarrow u\bar{u}s\bar{s}c\bar{c} + (n-6)g$	N/A	N/A	N/A
gen O3 one-flow, hel. sum	N/A	N/A	N/A
	N/A	N/A	N/A
gen O3 all-flows, one-hel	N/A	N/A	N/A
	N/A	N/A	N/A
run O3 one-flow, hel. sum	N/A	N/A	N/A
	N/A	N/A	N/A
run O3 all-flows, one-hel	N/A	N/A	N/A
	N/A	N/A	N/A

Built-in SM LC process matrix (continued)

ID base process	n=7	n=8	n=9
1 $d\bar{d} \rightarrow Z + (n-1)g$	N/A	N/A	N/A
2 $u\bar{d} \rightarrow W^+ + (n-1)g$	N/A	N/A	N/A
3 $d\bar{d} \rightarrow e^+e^- + (n-2)g$	N/A	N/A	N/A
4 $u\bar{d} \rightarrow e^+\nu_e + (n-2)g$	N/A	N/A	N/A
5 $d\bar{d} \rightarrow ZZ + (n-2)g$	N/A	N/A	N/A
6 $gg \rightarrow t\bar{t} + (n-2)g$	N/A	N/A	N/A
7 $d\bar{d} \rightarrow t\bar{t} + (n-2)g$	N/A	N/A	N/A
8 $gg \rightarrow gg + (n-2)g$	N/A	N/A	N/A
9 $d\bar{d} \rightarrow ZZZ + (n-3)g$	N/A	N/A	N/A
10 $d\bar{d} \rightarrow e^+e^-ZH + (n-4)g$	N/A	N/A	N/A
11 $d\bar{d} \rightarrow t\bar{t}ZH + (n-4)g$	N/A	N/A	N/A
12 $d\bar{d} \rightarrow e^+e^-e^+e^- + (n-4)g$	N/A	N/A	N/A
13 $d\bar{d} \rightarrow u\bar{u}s\bar{s} + (n-4)g$	N/A	N/A	N/A
14 $d\bar{d} \rightarrow u\bar{u}s\bar{s}c\bar{c} + (n-6)g$	N/A	N/A	N/A
gen O3 one-flow, hel. sum	N/A	N/A	N/A
	N/A	N/A	N/A
gen O3 all-flows, one-hel	N/A	N/A	N/A
	N/A	N/A	N/A
run O3 one-flow, hel. sum	N/A	N/A	N/A
	N/A	N/A	N/A
run O3 all-flows, one-hel	N/A	N/A	N/A
	N/A	N/A	N/A

Cell ordering and units follow the reading guide at the start of this section. A slash separates the two AmpliCol LC runtime workloads; a vertical bar separates the pyAmpliCol Rusticol-core wall and evaluator-only multipliers. Every multiplier is pyAmpliCol divided by AmpliCol.

5.2 Built-in SM NLC process matrix

ID	base process	n=1	n=2	n=3
1	$d\bar{d} \rightarrow Z + (n-1)g$	N/A	N/A	N/A
2	$u\bar{d} \rightarrow W^+ + (n-1)g$	N/A	N/A	N/A
3	$d\bar{d} \rightarrow e^+e^- + (n-2)g$	N/A	N/A	N/A
4	$u\bar{d} \rightarrow e^+\nu_e + (n-2)g$	N/A	N/A	N/A
5	$d\bar{d} \rightarrow ZZ + (n-2)g$	N/A	N/A	N/A
6	$gg \rightarrow t\bar{t} + (n-2)g$	N/A	N/A	N/A
7	$d\bar{d} \rightarrow t\bar{t} + (n-2)g$	N/A	N/A	N/A
8	$gg \rightarrow gg + (n-2)g$	N/A	N/A	N/A
9	$d\bar{d} \rightarrow ZZZ + (n-3)g$	N/A	N/A	N/A
10	$d\bar{d} \rightarrow e^+e^-ZH + (n-4)g$	N/A	N/A	N/A
11	$d\bar{d} \rightarrow t\bar{t}ZH + (n-4)g$	N/A	N/A	N/A
12	$d\bar{d} \rightarrow e^+e^-e^+e^- + (n-4)g$	N/A	N/A	N/A
13	$d\bar{d} \rightarrow u\bar{u}s\bar{s} + (n-4)g$	N/A	N/A	N/A
14	$d\bar{d} \rightarrow u\bar{u}s\bar{s}c\bar{c} + (n-6)g$	N/A	N/A	N/A
summary: gen		N/A	N/A	N/A
summary: run wall		N/A	N/A	N/A
summary: run eval		N/A	N/A	N/A

Built-in SM NLC process matrix (continued)

ID	base process	n=4	n=5
1	$d\bar{d} \rightarrow Z + (n-1)g$	N/A	N/A
2	$u\bar{d} \rightarrow W^+ + (n-1)g$	N/A	N/A
3	$d\bar{d} \rightarrow e^+e^- + (n-2)g$	N/A	N/A
4	$u\bar{d} \rightarrow e^+\nu_e + (n-2)g$	N/A	N/A
5	$d\bar{d} \rightarrow ZZ + (n-2)g$	N/A	N/A
6	$gg \rightarrow t\bar{t} + (n-2)g$	N/A	N/A
7	$d\bar{d} \rightarrow t\bar{t} + (n-2)g$	N/A	N/A
8	$gg \rightarrow gg + (n-2)g$	N/A	N/A
9	$d\bar{d} \rightarrow ZZZ + (n-3)g$	N/A	N/A
10	$d\bar{d} \rightarrow e^+e^-ZH + (n-4)g$	N/A	N/A
11	$d\bar{d} \rightarrow t\bar{t}ZH + (n-4)g$	N/A	N/A
12	$d\bar{d} \rightarrow e^+e^-e^+e^- + (n-4)g$	N/A	N/A
13	$d\bar{d} \rightarrow u\bar{u}s\bar{s} + (n-4)g$	N/A	N/A
14	$d\bar{d} \rightarrow u\bar{u}s\bar{s}c\bar{c} + (n-6)g$	N/A	N/A
summary: gen		N/A	N/A
summary: run wall		N/A	N/A
summary: run eval		N/A	N/A

Cell ordering and units follow the reading guide at the start of this section. A slash separates the two AmpliCol LC runtime workloads; a vertical bar separates the pyAmpliCol Rusticol-core wall and evaluator-only multipliers. Every multiplier is pyAmpliCol divided by AmpliCol.

5.3 Built-in SM full-colour process matrix

ID	base process	n=1	n=2	n=3
1	$d\bar{d} \rightarrow Z + (n-1)g$	N/A	N/A	N/A
2	$u\bar{d} \rightarrow W^+ + (n-1)g$	N/A	N/A	N/A
3	$d\bar{d} \rightarrow e^+e^- + (n-2)g$	N/A	N/A	N/A
4	$u\bar{d} \rightarrow e^+\nu_e + (n-2)g$	N/A	N/A	N/A
5	$d\bar{d} \rightarrow ZZ + (n-2)g$	N/A	N/A	N/A
6	$gg \rightarrow t\bar{t} + (n-2)g$	N/A	N/A	N/A
7	$d\bar{d} \rightarrow t\bar{t} + (n-2)g$	N/A	N/A	N/A
8	$gg \rightarrow gg + (n-2)g$	N/A	N/A	N/A
9	$d\bar{d} \rightarrow ZZZ + (n-3)g$	N/A	N/A	N/A
10	$d\bar{d} \rightarrow e^+e^-ZH + (n-4)g$	N/A	N/A	N/A
11	$d\bar{d} \rightarrow t\bar{t}ZH + (n-4)g$	N/A	N/A	N/A
12	$d\bar{d} \rightarrow e^+e^-e^+e^- + (n-4)g$	N/A	N/A	N/A
13	$d\bar{d} \rightarrow u\bar{u}s\bar{s} + (n-4)g$	N/A	N/A	N/A
14	$d\bar{d} \rightarrow u\bar{u}s\bar{s}c\bar{c} + (n-6)g$	N/A	N/A	N/A
summary: gen		N/A	N/A	N/A
summary: run wall		N/A	N/A	N/A
summary: run eval		N/A	N/A	N/A

Built-in SM full-colour process matrix (continued)

ID	base process	n=4	n=5
1	$d\bar{d} \rightarrow Z + (n-1)g$	N/A	N/A
2	$u\bar{d} \rightarrow W^+ + (n-1)g$	N/A	N/A
3	$d\bar{d} \rightarrow e^+e^- + (n-2)g$	N/A	N/A
4	$u\bar{d} \rightarrow e^+\nu_e + (n-2)g$	N/A	N/A
5	$d\bar{d} \rightarrow ZZ + (n-2)g$	N/A	N/A
6	$gg \rightarrow t\bar{t} + (n-2)g$	N/A	N/A
7	$d\bar{d} \rightarrow t\bar{t} + (n-2)g$	N/A	N/A
8	$gg \rightarrow gg + (n-2)g$	N/A	N/A
9	$d\bar{d} \rightarrow ZZZ + (n-3)g$	N/A	N/A
10	$d\bar{d} \rightarrow e^+e^-ZH + (n-4)g$	N/A	N/A
11	$d\bar{d} \rightarrow t\bar{t}ZH + (n-4)g$	N/A	N/A
12	$d\bar{d} \rightarrow e^+e^-e^+e^- + (n-4)g$	N/A	N/A
13	$d\bar{d} \rightarrow u\bar{u}s\bar{s} + (n-4)g$	N/A	N/A
14	$d\bar{d} \rightarrow u\bar{u}s\bar{s}c\bar{c} + (n-6)g$	N/A	N/A
summary: gen		N/A	N/A
summary: run wall		N/A	N/A
summary: run eval		N/A	N/A

Cell ordering and units follow the reading guide at the start of this section. A slash separates the two AmpliCol LC runtime workloads; a vertical bar separates the pyAmpliCol Rusticol-core wall and evaluator-only multipliers. Every multiplier is pyAmpliCol divided by AmpliCol.

5.4 UFO-SM LC process matrix

ID	base process	n=1	n=2	n=3
1	$d\bar{d} \rightarrow Z + (n-1)g$	N/A	N/A	N/A
2	$u\bar{d} \rightarrow W^+ + (n-1)g$	N/A	N/A	N/A
3	$d\bar{d} \rightarrow e^+e^- + (n-2)g$	N/A	N/A	N/A
4	$u\bar{d} \rightarrow e^+\nu_e + (n-2)g$	N/A	N/A	N/A
5	$d\bar{d} \rightarrow ZZ + (n-2)g$	N/A	N/A	N/A
6	$gg \rightarrow t\bar{t} + (n-2)g$	N/A	N/A	N/A
7	$d\bar{d} \rightarrow t\bar{t} + (n-2)g$	N/A	N/A	N/A
8	$gg \rightarrow gg + (n-2)g$	N/A	N/A	N/A
9	$d\bar{d} \rightarrow ZZZ + (n-3)g$	N/A	N/A	N/A
10	$d\bar{d} \rightarrow e^+e^-ZH + (n-4)g$	N/A	N/A	N/A
11	$d\bar{d} \rightarrow t\bar{t}ZH + (n-4)g$	N/A	N/A	N/A
12	$d\bar{d} \rightarrow e^+e^-e^+e^- + (n-4)g$	N/A	N/A	N/A
13	$d\bar{d} \rightarrow u\bar{u}s\bar{s} + (n-4)g$	N/A	N/A	N/A
14	$d\bar{d} \rightarrow u\bar{u}s\bar{s}c\bar{c} + (n-6)g$	N/A	N/A	N/A
gen O3 one-flow, hel. sum		N/A	N/A	N/A
gen O3 all-flows, one-hel		N/A	N/A	N/A
run O3 one-flow, hel. sum		N/A	N/A	N/A
run O3 all-flows, one-hel		N/A	N/A	N/A

UFO-SM LC process matrix (continued)

ID	base process	n=4	n=5	n=6
1	$d\bar{d} \rightarrow Z + (n-1)g$	N/A	N/A	N/A
2	$u\bar{d} \rightarrow W^+ + (n-1)g$	N/A	N/A	N/A
3	$d\bar{d} \rightarrow e^+e^- + (n-2)g$	N/A	N/A	N/A
4	$u\bar{d} \rightarrow e^+\nu_e + (n-2)g$	N/A	N/A	N/A
5	$d\bar{d} \rightarrow ZZ + (n-2)g$	N/A	N/A	N/A
6	$gg \rightarrow t\bar{t} + (n-2)g$	N/A	N/A	N/A
7	$d\bar{d} \rightarrow t\bar{t} + (n-2)g$	N/A	N/A	N/A
8	$gg \rightarrow gg + (n-2)g$	N/A	N/A	N/A
9	$d\bar{d} \rightarrow ZZZ + (n-3)g$	N/A	N/A	N/A
10	$d\bar{d} \rightarrow e^+e^-ZH + (n-4)g$	N/A	N/A	N/A
11	$d\bar{d} \rightarrow t\bar{t}ZH + (n-4)g$	N/A	N/A	N/A
12	$d\bar{d} \rightarrow e^+e^-e^+e^- + (n-4)g$	N/A	N/A	N/A
13	$d\bar{d} \rightarrow u\bar{u}s\bar{s} + (n-4)g$	N/A	N/A	N/A
14	$d\bar{d} \rightarrow u\bar{u}s\bar{s}c\bar{c} + (n-6)g$	N/A	N/A	N/A
gen O3 one-flow, hel. sum		N/A	N/A	N/A
gen O3 all-flows, one-hel		N/A	N/A	N/A
run O3 one-flow, hel. sum		N/A	N/A	N/A
run O3 all-flows, one-hel		N/A	N/A	N/A

UFO-SM LC process matrix (continued)

ID	base process	n=7	n=8	n=9
1	$d\bar{d} \rightarrow Z + (n-1)g$	N/A	N/A	N/A
2	$u\bar{d} \rightarrow W^+ + (n-1)g$	N/A	N/A	N/A
3	$d\bar{d} \rightarrow e^+e^- + (n-2)g$	N/A	N/A	N/A
4	$u\bar{d} \rightarrow e^+\nu_e + (n-2)g$	N/A	N/A	N/A
5	$d\bar{d} \rightarrow ZZ + (n-2)g$	N/A	N/A	N/A
6	$gg \rightarrow t\bar{t} + (n-2)g$	N/A	N/A	N/A
7	$d\bar{d} \rightarrow t\bar{t} + (n-2)g$	N/A	N/A	N/A
8	$gg \rightarrow gg + (n-2)g$	N/A	N/A	N/A
9	$d\bar{d} \rightarrow ZZZ + (n-3)g$	N/A	N/A	N/A
10	$d\bar{d} \rightarrow e^+e^-ZH + (n-4)g$	N/A	N/A	N/A
11	$d\bar{d} \rightarrow t\bar{t}ZH + (n-4)g$	N/A	N/A	N/A
12	$d\bar{d} \rightarrow e^+e^-e^+e^- + (n-4)g$	N/A	N/A	N/A
13	$d\bar{d} \rightarrow u\bar{u}s\bar{s} + (n-4)g$	N/A	N/A	N/A
14	$d\bar{d} \rightarrow u\bar{u}s\bar{s}c\bar{c} + (n-6)g$	N/A	N/A	N/A
gen O3 one-flow, hel. sum		N/A	N/A	N/A
gen O3 all-flows, one-hel		N/A	N/A	N/A
run O3 one-flow, hel. sum		N/A	N/A	N/A
run O3 all-flows, one-hel		N/A	N/A	N/A

Cell ordering and units follow the reading guide at the start of this section. A slash separates the two AmpliCol LC runtime workloads; a vertical bar separates the pyAmpliCol Rusticol-core wall and evaluator-only multipliers. Every multiplier is pyAmpliCol divided by AmpliCol.

5.5 UFO-SM NLC process matrix

ID	base process	n=1	n=2	n=3
1	$d\bar{d} \rightarrow Z + (n-1)g$	N/A	N/A	N/A
2	$u\bar{d} \rightarrow W^+ + (n-1)g$	N/A	N/A	N/A
3	$d\bar{d} \rightarrow e^+e^- + (n-2)g$	N/A	N/A	N/A
4	$u\bar{d} \rightarrow e^+\nu_e + (n-2)g$	N/A	N/A	N/A
5	$d\bar{d} \rightarrow ZZ + (n-2)g$	N/A	N/A	N/A
6	$gg \rightarrow t\bar{t} + (n-2)g$	N/A	N/A	N/A
7	$d\bar{d} \rightarrow t\bar{t} + (n-2)g$	N/A	N/A	N/A
8	$gg \rightarrow gg + (n-2)g$	N/A	N/A	N/A
9	$d\bar{d} \rightarrow ZZZ + (n-3)g$	N/A	N/A	N/A
10	$d\bar{d} \rightarrow e^+e^-ZH + (n-4)g$	N/A	N/A	N/A
11	$d\bar{d} \rightarrow t\bar{t}ZH + (n-4)g$	N/A	N/A	N/A
12	$d\bar{d} \rightarrow e^+e^-e^+e^- + (n-4)g$	N/A	N/A	N/A
13	$d\bar{d} \rightarrow u\bar{u}s\bar{s} + (n-4)g$	N/A	N/A	N/A
14	$d\bar{d} \rightarrow u\bar{u}s\bar{s}c\bar{c} + (n-6)g$	N/A	N/A	N/A
summary: gen		N/A	N/A	N/A
summary: run wall		N/A	N/A	N/A
summary: run eval		N/A	N/A	N/A

UFO-SM NLC process matrix (continued)

ID	base process	n=4	n=5
1	$d\bar{d} \rightarrow Z + (n-1)g$	N/A	N/A
2	$u\bar{d} \rightarrow W^+ + (n-1)g$	N/A	N/A
3	$d\bar{d} \rightarrow e^+e^- + (n-2)g$	N/A	N/A
4	$u\bar{d} \rightarrow e^+\nu_e + (n-2)g$	N/A	N/A
5	$d\bar{d} \rightarrow ZZ + (n-2)g$	N/A	N/A
6	$gg \rightarrow t\bar{t} + (n-2)g$	N/A	N/A
7	$d\bar{d} \rightarrow t\bar{t} + (n-2)g$	N/A	N/A
8	$gg \rightarrow gg + (n-2)g$	N/A	N/A
9	$d\bar{d} \rightarrow ZZZ + (n-3)g$	N/A	N/A
10	$d\bar{d} \rightarrow e^+e^-ZH + (n-4)g$	N/A	N/A
11	$d\bar{d} \rightarrow t\bar{t}ZH + (n-4)g$	N/A	N/A
12	$d\bar{d} \rightarrow e^+e^-e^+e^- + (n-4)g$	N/A	N/A
13	$d\bar{d} \rightarrow u\bar{u}s\bar{s} + (n-4)g$	N/A	N/A
14	$d\bar{d} \rightarrow u\bar{u}s\bar{s}c\bar{c} + (n-6)g$	N/A	N/A
summary: gen		N/A	N/A
summary: run wall		N/A	N/A
summary: run eval		N/A	N/A

Cell ordering and units follow the reading guide at the start of this section. A slash separates the two AmpliCol LC runtime workloads; a vertical bar separates the pyAmpliCol Rusticol-core wall and evaluator-only multipliers. Every multiplier is pyAmpliCol divided by AmpliCol.

5.6 UFO-SM full-colour process matrix

ID	base process	n=1	n=2	n=3
1	$d\bar{d} \rightarrow Z + (n-1)g$	N/A	N/A	N/A
2	$u\bar{d} \rightarrow W^+ + (n-1)g$	N/A	N/A	N/A
3	$d\bar{d} \rightarrow e^+e^- + (n-2)g$	N/A	N/A	N/A
4	$u\bar{d} \rightarrow e^+\nu_e + (n-2)g$	N/A	N/A	N/A
5	$d\bar{d} \rightarrow ZZ + (n-2)g$	N/A	N/A	N/A
6	$gg \rightarrow t\bar{t} + (n-2)g$	N/A	N/A	N/A
7	$d\bar{d} \rightarrow t\bar{t} + (n-2)g$	N/A	N/A	N/A
8	$gg \rightarrow gg + (n-2)g$	N/A	N/A	N/A
9	$d\bar{d} \rightarrow ZZZ + (n-3)g$	N/A	N/A	N/A
10	$d\bar{d} \rightarrow e^+e^-ZH + (n-4)g$	N/A	N/A	N/A
11	$d\bar{d} \rightarrow t\bar{t}ZH + (n-4)g$	N/A	N/A	N/A
12	$d\bar{d} \rightarrow e^+e^-e^+e^- + (n-4)g$	N/A	N/A	N/A
13	$d\bar{d} \rightarrow u\bar{u}s\bar{s} + (n-4)g$	N/A	N/A	N/A
14	$d\bar{d} \rightarrow u\bar{u}s\bar{s}c\bar{c} + (n-6)g$	N/A	N/A	N/A
summary: gen		N/A	N/A	N/A
summary: run wall		N/A	N/A	N/A
summary: run eval		N/A	N/A	N/A

UFO-SM full-colour process matrix (continued)

ID	base process	n=4	n=5
1	$d\bar{d} \rightarrow Z + (n-1)g$	N/A	N/A
2	$u\bar{d} \rightarrow W^+ + (n-1)g$	N/A	N/A
3	$d\bar{d} \rightarrow e^+e^- + (n-2)g$	N/A	N/A
4	$u\bar{d} \rightarrow e^+\nu_e + (n-2)g$	N/A	N/A
5	$d\bar{d} \rightarrow ZZ + (n-2)g$	N/A	N/A
6	$gg \rightarrow t\bar{t} + (n-2)g$	N/A	N/A
7	$d\bar{d} \rightarrow t\bar{t} + (n-2)g$	N/A	N/A
8	$gg \rightarrow gg + (n-2)g$	N/A	N/A
9	$d\bar{d} \rightarrow ZZZ + (n-3)g$	N/A	N/A
10	$d\bar{d} \rightarrow e^+e^-ZH + (n-4)g$	N/A	N/A
11	$d\bar{d} \rightarrow t\bar{t}ZH + (n-4)g$	N/A	N/A
12	$d\bar{d} \rightarrow e^+e^-e^+e^- + (n-4)g$	N/A	N/A
13	$d\bar{d} \rightarrow u\bar{u}s\bar{s} + (n-4)g$	N/A	N/A
14	$d\bar{d} \rightarrow u\bar{u}s\bar{s}c\bar{c} + (n-6)g$	N/A	N/A
summary: gen		N/A	N/A
summary: run wall		N/A	N/A
summary: run eval		N/A	N/A

Cell ordering and units follow the reading guide at the start of this section. A slash separates the two AmpliCol LC runtime workloads; a vertical bar separates the pyAmpliCol Rusticol-core wall and evaluator-only multipliers. Every multiplier is pyAmpliCol divided by AmpliCol.

6 Z-plus-Jets Evaluator Ladders

The dedicated $d\bar{d} \rightarrow Z + (n-1)g$ ladders compare the independent reference with JIT level 1, JIT level 3, ASM O3, and C++ O3 evaluator configurations. They retain final-state multiplicities $n = 1, \dots, 9$ for the built-in model and $n = 1, \dots, 6$ for the UFO-SM campaign. Each row states its evaluator backend and optimization level explicitly; all other workload settings follow the benchmark contract above.

Runtime comparisons between these backends also include their present vectorization strategy. The SymJIT application evaluates a batch with SIMD across phase-space points, whereas the current generated ASM and C++ payloads are scalar and Rusticol invokes them point by point. Their ratios therefore measure both symbolic/code-generation quality and this SIMD advantage; they must not be interpreted as a comparison of otherwise identical vectorized kernels.

Worked $d\bar{d} \rightarrow Zg$ example

For $n = 2$, the explicit process is

$$d(p_1) \bar{d}(p_2) \rightarrow Z(p_3) g(p_4).$$

The report labels source particles in process order. The selected leading-colour reference order for the selected-flow timing is therefore (2, 4, 1, 3): the colour line enters through $\bar{d}$, traverses the gluon, and exits through d , while the colour-singlet Z remains in the public ordering only to keep source labels unambiguous. Internally, the colour-sector match is the coloured word (2, 4, 1).

`pyAmpliCol` builds one shared-current DAG for the selected process. Source nodes carry particle, momentum, helicity, flavour, and colour-flow labels; composite nodes are keyed by the full current identity and the external-label mask. When two diagrams need the same current, the DAG contains one node and downstream amplitude roots reuse its value. In the selected-flow/helicity-sum workload the generator keeps the colour sector compatible with (2, 4, 1) and sums source helicities. In the all-flows/fixed-helicity workload it keeps all LC colour flows but fixes one DAG-supported source-helicity assignment. The generated artifact serializes the stages and evaluator payloads; Rusticol loads that artifact, executes `Runtime.evaluate()` inside the native core, and `BenchmarkRunner` records both the native repeated wall time and the sum of evaluator-call timings.

6.1 Dedicated $d\bar{d} \rightarrow Z +$ Gluon Performance

n process	route	setup	selected flow, helicity sum			all flows, fixed helicity		
			gen [s]	wall [us/pt]	eval [us/pt]	gen [s]	wall [us/pt]	eval [us/pt]
1 d d ⁻ > Z + (1-1)*g	reference	AmpliCol	N/A	N/A	N/I	N/A	N/A	N/A
1 d d ⁻ > Z + (1-1)*g	Rusticol	pyAmpliCol JIT O1	N/A	N/A	N/A	N/A	N/A	N/A
1 d d ⁻ > Z + (1-1)*g	Rusticol	pyAmpliCol ASM O3	N/A	N/A	N/A	N/A	N/A	N/A
1 d d ⁻ > Z + (1-1)*g	Rusticol	pyAmpliCol C++ O3	N/A	N/A	N/A	N/A	N/A	N/A
1 d d ⁻ > Z + (1-1)*g	Rusticol	pyAmpliCol JIT O3	N/A	N/A	N/I	N/A	N/A	N/A
2 d d ⁻ > Z + (2-1)*g	reference	AmpliCol	N/A	N/A	N/I	N/A	N/A	N/A
2 d d ⁻ > Z + (2-1)*g	Rusticol	pyAmpliCol JIT O1	N/A	N/A	N/A	N/A	N/A	N/A
2 d d ⁻ > Z + (2-1)*g	Rusticol	pyAmpliCol ASM O3	N/A	N/A	N/A	N/A	N/A	N/A
2 d d ⁻ > Z + (2-1)*g	Rusticol	pyAmpliCol C++ O3	N/A	N/A	N/A	N/A	N/A	N/A
2 d d ⁻ > Z + (2-1)*g	Rusticol	pyAmpliCol JIT O3	N/A	N/A	N/I	N/A	N/A	N/A
3 d d ⁻ > Z + (3-1)*g	reference	AmpliCol	N/A	N/A	N/I	N/A	N/A	N/A
3 d d ⁻ > Z + (3-1)*g	Rusticol	pyAmpliCol JIT O1	N/A	N/A	N/A	N/A	N/A	N/A
3 d d ⁻ > Z + (3-1)*g	Rusticol	pyAmpliCol ASM O3	N/A	N/A	N/A	N/A	N/A	N/A
3 d d ⁻ > Z + (3-1)*g	Rusticol	pyAmpliCol C++ O3	N/A	N/A	N/A	N/A	N/A	N/A
3 d d ⁻ > Z + (3-1)*g	Rusticol	pyAmpliCol JIT O3	N/A	N/A	N/I	N/A	N/A	N/A
4 d d ⁻ > Z + (4-1)*g	reference	AmpliCol	N/A	N/A	N/I	N/A	N/A	N/A
4 d d ⁻ > Z + (4-1)*g	Rusticol	pyAmpliCol JIT O1	N/A	N/A	N/A	N/A	N/A	N/A
4 d d ⁻ > Z + (4-1)*g	Rusticol	pyAmpliCol ASM O3	N/A	N/A	N/A	N/A	N/A	N/A
4 d d ⁻ > Z + (4-1)*g	Rusticol	pyAmpliCol C++ O3	N/A	N/A	N/A	N/A	N/A	N/A

n process	route	setup	selected flow, helicity sum			all flows, fixed helicity		
			gen [s]	wall [us/pt]	eval [us/pt]	gen [s]	wall [us/pt]	eval [us/pt]
4 d d ⁻ > Z + (4-1)*g	Rusticol	pyAmpliCol JIT O3	N/A	N/A	N/I	N/A	N/A	N/A
5 d d ⁻ > Z + (5-1)*g	reference	AmpliCol	N/A	N/A	N/I	N/A	N/A	N/A
5 d d ⁻ > Z + (5-1)*g	Rusticol	pyAmpliCol JIT O1	N/A	N/A	N/A	N/A	N/A	N/A
5 d d ⁻ > Z + (5-1)*g	Rusticol	pyAmpliCol ASM O3	N/A	N/A	N/A	N/A	N/A	N/A
5 d d ⁻ > Z + (5-1)*g	Rusticol	pyAmpliCol C++ O3	N/A	N/A	N/A	N/A	N/A	N/A
5 d d ⁻ > Z + (5-1)*g	Rusticol	pyAmpliCol JIT O3	N/A	N/A	N/I	N/A	N/A	N/A
6 d d ⁻ > Z + (6-1)*g	reference	AmpliCol	N/A	N/A	N/I	N/A	N/A	N/A
6 d d ⁻ > Z + (6-1)*g	Rusticol	pyAmpliCol JIT O1	N/A	N/A	N/A	N/A	N/A	N/A
6 d d ⁻ > Z + (6-1)*g	Rusticol	pyAmpliCol ASM O3	N/A	N/A	N/A	N/A	N/A	N/A
6 d d ⁻ > Z + (6-1)*g	Rusticol	pyAmpliCol C++ O3	N/A	N/A	N/A	N/A	N/A	N/A
6 d d ⁻ > Z + (6-1)*g	Rusticol	pyAmpliCol JIT O3	N/A	N/A	N/I	N/A	N/A	N/A
7 d d ⁻ > Z + (7-1)*g	reference	AmpliCol	N/A	N/A	N/I	N/A	N/A	N/A
7 d d ⁻ > Z + (7-1)*g	Rusticol	pyAmpliCol JIT O1	N/A	N/A	N/A	N/A	N/A	N/A
7 d d ⁻ > Z + (7-1)*g	Rusticol	pyAmpliCol ASM O3	N/A	N/A	N/A	N/A	N/A	N/A
7 d d ⁻ > Z + (7-1)*g	Rusticol	pyAmpliCol C++ O3	N/A	N/A	N/A	N/A	N/A	N/A
7 d d ⁻ > Z + (7-1)*g	Rusticol	pyAmpliCol JIT O3	N/A	N/A	N/I	N/A	N/A	N/A
8 d d ⁻ > Z + (8-1)*g	reference	AmpliCol	N/A	N/A	N/I	N/A	N/A	N/A
8 d d ⁻ > Z + (8-1)*g	Rusticol	pyAmpliCol JIT O1	N/A	N/A	N/A	N/A	N/A	N/A
8 d d ⁻ > Z + (8-1)*g	Rusticol	pyAmpliCol ASM O3	N/A	N/A	N/A	N/A	N/A	N/A
8 d d ⁻ > Z + (8-1)*g	Rusticol	pyAmpliCol C++ O3	N/A	N/A	N/A	N/A	N/A	N/A
8 d d ⁻ > Z + (8-1)*g	Rusticol	pyAmpliCol JIT O3	N/A	N/A	N/I	N/A	N/A	N/A
9 d d ⁻ > Z + (9-1)*g	reference	AmpliCol	N/A	N/A	N/I	N/A	N/A	N/A
9 d d ⁻ > Z + (9-1)*g	Rusticol	pyAmpliCol JIT O1	N/A	N/A	N/A	N/A	N/A	N/A
9 d d ⁻ > Z + (9-1)*g	Rusticol	pyAmpliCol ASM O3	N/A	N/A	N/A	N/A	N/A	N/A
9 d d ⁻ > Z + (9-1)*g	Rusticol	pyAmpliCol C++ O3	N/A	N/A	N/A	N/A	N/A	N/A
9 d d ⁻ > Z + (9-1)*g	Rusticol	pyAmpliCol JIT O3	N/A	N/A	N/I	N/A	N/A	N/A

6.2 UFO-SM Dedicated $d\bar{d} \rightarrow Z + \text{Gluon}$ Performance

n process	route	setup	selected flow, helicity sum				all flows, fixed helicity					
			gen [s]	vs blt-in	wall [us/pt]	vs blt-in	eval [us/pt]	gen [s]	vs blt-in	wall [us/pt]	vs blt-in	eval [us/pt]
1 d d ⁺ > Z + (1-1)*g	reference	AmpliCol	N/A		N/A		N/I	N/A		N/A		N/A
1 d d ⁺ > Z + (1-1)*g	Rusticol	pyAmpliCol JIT O1	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A
1 d d ⁺ > Z + (1-1)*g	Rusticol	pyAmpliCol ASM O3	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A
1 d d ⁺ > Z + (1-1)*g	Rusticol	pyAmpliCol C++ O3	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A
1 d d ⁺ > Z + (1-1)*g	Rusticol	pyAmpliCol JIT O3	N/A	N/A	N/A	N/A	N/I	N/A	N/A	N/A	N/A	N/A
2 d d ⁺ > Z + (2-1)*g	reference	AmpliCol	N/A		N/A		N/I	N/A		N/A		N/A
2 d d ⁺ > Z + (2-1)*g	Rusticol	pyAmpliCol JIT O1	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A
2 d d ⁺ > Z + (2-1)*g	Rusticol	pyAmpliCol ASM O3	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A
2 d d ⁺ > Z + (2-1)*g	Rusticol	pyAmpliCol C++ O3	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A
2 d d ⁺ > Z + (2-1)*g	Rusticol	pyAmpliCol JIT O3	N/A	N/A	N/A	N/A	N/I	N/A	N/A	N/A	N/A	N/A
3 d d ⁺ > Z + (3-1)*g	reference	AmpliCol	N/A		N/A		N/I	N/A		N/A		N/A
3 d d ⁺ > Z + (3-1)*g	Rusticol	pyAmpliCol JIT O1	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A
3 d d ⁺ > Z + (3-1)*g	Rusticol	pyAmpliCol ASM O3	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A
3 d d ⁺ > Z + (3-1)*g	Rusticol	pyAmpliCol C++ O3	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A
3 d d ⁺ > Z + (3-1)*g	Rusticol	pyAmpliCol JIT O3	N/A	N/A	N/A	N/A	N/I	N/A	N/A	N/A	N/A	N/A
4 d d ⁺ > Z + (4-1)*g	reference	AmpliCol	N/A		N/A		N/I	N/A		N/A		N/A
4 d d ⁺ > Z + (4-1)*g	Rusticol	pyAmpliCol JIT O1	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A
4 d d ⁺ > Z + (4-1)*g	Rusticol	pyAmpliCol ASM O3	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A
4 d d ⁺ > Z + (4-1)*g	Rusticol	pyAmpliCol C++ O3	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A
4 d d ⁺ > Z + (4-1)*g	Rusticol	pyAmpliCol JIT O3	N/A	N/A	N/A	N/A	N/I	N/A	N/A	N/A	N/A	N/A
5 d d ⁺ > Z + (5-1)*g	reference	AmpliCol	N/A		N/A		N/I	N/A		N/A		N/A
5 d d ⁺ > Z + (5-1)*g	Rusticol	pyAmpliCol JIT O1	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A
5 d d ⁺ > Z + (5-1)*g	Rusticol	pyAmpliCol ASM O3	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A
5 d d ⁺ > Z + (5-1)*g	Rusticol	pyAmpliCol C++ O3	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A
5 d d ⁺ > Z + (5-1)*g	Rusticol	pyAmpliCol JIT O3	N/A	N/A	N/A	N/A	N/I	N/A	N/A	N/A	N/A	N/A
6 d d ⁺ > Z + (6-1)*g	reference	AmpliCol	N/A		N/A		N/I	N/A		N/A		N/A
6 d d ⁺ > Z + (6-1)*g	Rusticol	pyAmpliCol JIT O1	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A
6 d d ⁺ > Z + (6-1)*g	Rusticol	pyAmpliCol ASM O3	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A
6 d d ⁺ > Z + (6-1)*g	Rusticol	pyAmpliCol C++ O3	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A
6 d d ⁺ > Z + (6-1)*g	Rusticol	pyAmpliCol JIT O3	N/A	N/A	N/A	N/A	N/I	N/A	N/A	N/A	N/A	N/A

Each workload block reports generation in seconds and Rusticol-core wall and evaluator-only runtime in microseconds per point. Parenthesized multipliers compare pyAmpliCol with AmpliCol; the UFO-SM “vs blt-in” columns compare the same UFO-SM and built-in-SM setup.

7 Colour-Singlet Model Ladders

The scalar-contact ladder retains $n = 2, \dots, 8$ final scalars. It tests generic high-valence contact decomposition, identical-particle normalization, and source parity without Standard-Model-specific kernels. The scalar-gravity ladder retains $n = 2, \dots, 4$ final gravitons and exercises spin-two wavefunctions, tensor propagators, and momentum-dependent vertices. Since all external particles in these fixtures are colour singlets, LC, NLC, and full colour reduce to the same unit colour contraction.

For these ladders, the relative-difference row is $|M_{f64} - M_{hp}| / \max(|M_{hp}|, 10^{-300})$, where M_{hp} is the same point re-evaluated through the higher-precision diagnostic path. It is a numerical stability check, not a performance ratio.

7.1 Scalar-contact ladder

metric	scalar_0 scalar_0 > n*scalar_0						
	n = 2	n = 3	n = 4	n = 5	n = 6	n = 7	n = 8
generation [s]	N/A	N/A	N/A	N/A	N/A	N/A	N/A
wall [μ s/pt]	N/A	N/A	N/A	N/A	N/A	N/A	N/A
evaluator [μ s/pt]	N/A	N/A	N/A	N/A	N/A	N/A	N/A
matrix element	N/A	N/A	N/A	N/A	N/A	N/A	N/A
relative diff. vs hp	N/A	N/A	N/A	N/A	N/A	N/A	N/A

7.2 Scalar-gravity ladder

metric	scalar_0 scalar_0 > n*graviton		
	$n = 2$	$n = 3$	$n = 4$
generation [s]	N/A	N/A	N/A
wall [μ s/pt]	N/A	N/A	N/A
evaluator [μ s/pt]	N/A	N/A	N/A
matrix element	N/A	N/A	N/A
relative diff. vs hp	N/A	N/A	N/A

8 Interpretation

This report establishes coverage, methodology, and the current measured result set. Performance comparisons are meaningful only for a reviewed baseline produced from the same effective configuration on named hardware. Readers should treat every N/A entry as unavailable and should not interpolate, copy, or estimate a value from neighbouring multiplicities or model profiles.

The retained measurement records under `docs/results/` are authoritative for table completeness and provenance. The TeX inputs and this PDF are their publication presentation.

9 Worked $d\bar{d} \rightarrow Zgg$ Example

This section follows one process from its user-facing notation to the amplitude roots generated by `pyAmpliCol`. It is deliberately concrete: the counts and labels below are those produced by the current built-in Standard Model for complete leading-colour (LC) coverage, not a schematic diagram count. The same process is also contained in the UFO-SM validation set.

9.1 External legs and conventions

The physical process is

$$d(p_1) \bar{d}(p_2) \longrightarrow Z(p_3) g(p_4) g(p_5), \quad p_1 + p_2 = p_3 + p_4 + p_5. \quad (4)$$

The integers $1, \dots, 5$ are *source labels*. They are stable labels of external particles in process order; they are neither current IDs nor colour-sector IDs. The DAG compiler crosses incoming particles once and then uses the all-outgoing momenta

$$q_1 = -p_1, \quad q_2 = -p_2, \quad q_i = p_i \ (i = 3, 4, 5), \quad \sum_{i=1}^5 q_i = 0. \quad (5)$$

Runtime clients nevertheless supply momenta in the physical order and sign convention of Eq. (4); Rusticol performs the two incoming sign changes when it fills momentum slots.

label	physical leg	DAG particle	DAG momentum	states	colour role
1	incoming d	outgoing $\bar{d}$	$q_1 = -p_1$	2	$\bar{\mathbf{3}}$, colour-line end
2	incoming $\bar{d}$	outgoing d	$q_2 = -p_2$	2	$\mathbf{3}$, colour-line start
3	outgoing Z	outgoing Z	$q_3 = p_3$	3	singlet
4	outgoing g	outgoing g	$q_4 = p_4$	2	adjoint insertion
5	outgoing g	outgoing g	$q_5 = p_5$	2	adjoint insertion

There are $2 \times 2 \times 3 \times 2 \times 2 = 48$ physical helicity configurations. The physics metadata retains all 48 stable helicity IDs, for example `h:+1,-1,+0,-1,+1`. Twenty-four configurations are proven structural zeros by the chiral source and kernel constraints and are returned as explicit zeros by resolved evaluation; only the nonzero representatives need amplitude roots.

At LC the singlet label 3 is absent from the coloured word. Complete coverage has the two physical flows

$$c_0 = (2, 4, 5, 1), \quad c_1 = (2, 5, 4, 1), \quad (6)$$

published as the stable runtime IDs `flow:2,4,5,1` and `flow:2,5,4,1`. A user selects these strings. The small integer sector used while building a DAG is artifact-local implementation data and must not be used as a public flow identifier.

9.2 From diagrams to reusable currents

For a nonempty source subset I , define

$$q_I = \sum_{i \in I} q_i, \quad m(I) = \sum_{i \in I} 2^{i-1}. \quad (7)$$

Thus $m(\{2, 4, 5\}) = 26$. A source wavefunction is a one-leg current $J_i^a = W_{t_i}^a(q_i, h_i)$, where W denotes the appropriate spinor, polarization vector, or scalar wavefunction. Composite currents obey the binary recursion

$$J_{t,\Omega}^a(I) = P_{t \ b}^a(q_I) \sum_{\substack{I=A \dot{\cup} B \\ v \text{phantom} I A, B \neq \emptyset}} \sum_{v: t_A t_B \rightarrow t} \kappa_v V_v^b{}_{cd}(q_A, q_B) J_{t_A, \Omega_A}^c(A) J_{t_B, \Omega_B}^d(B). \quad (8)$$

Here t is the particle state and Ω is the complete current identity described in Sec. 10. The sum contains only model, quantum-number, perturbative-order, and colour-order admissible combinations. The propagator P_t belongs to the output value slot; an unpropagated form is retained as well when a later closure needs it.

For one fixed source-state ancestry, representative terms along the quark line are

$$J_d(2, 5) = P_d V_{dgd} [J_d(2), J_g(5)], \quad (9)$$

$$J_g(4, 5) = P_g V_{ggg} [J_g(4), J_g(5)], \quad (10)$$

$$J_d(2, 4, 5) = P_d \left(V_{dgd} [J_d(2, 5), J_g(4)] + V_{dgd} [J_d(2, 4), J_g(5)] + V_{dgd} [J_d(2), J_g(4, 5)] \right), \quad (11)$$

$$J_d(2, 3, 4, 5) = P_d \sum_{A \dot{\cup} B = \{2, 3, 4, 5\}} V [J(A), J(B)]. \quad (12)$$

Equation (12) includes all admissible positions of the Z emission and both gluon traversals. The notation suppresses chirality, helicity ancestry, couplings, and ordered-label data; those fields are never suppressed in the actual key. In particular, $J_d(2, 5, 4)$ and $J_d(2, 4, 5)$ are distinct ordered currents even though they carry the same unordered subset and momentum.

The final closure joins the endpoint current to source label 1:

$$\mathcal{A}_{h,c}^{(r)} = C_{aa}^{(r)} J_d^a(1) J_d^a(2, 3, 4, 5), \quad \mathcal{A}_{h,c} = \sum_{r \in G(h,c)} \mathcal{A}_{h,c}^{(r)}. \quad (13)$$

The contraction $C^{(r)}$, root coupling, colour weight, and coherent-group identifier are model-owned records, rather than special cases in Rusticol.

9.3 The concrete graph

The current compiler creates the following deterministic topology. IDs are dense diagnostic indices in this artifact, not stable public selectors.

layer	curr.	attach.	kernel eval.	representative content
sources	11	0	0	two Weyl states on each quark leg, three Z states, and two states on each gluon; current IDs 0–10
$ I = 2$	18	18	18	$J_d(2, 3)$, $J_d(2, 4)$, $J_d(2, 5)$, and $J_g(4, 5)$; current IDs 11–28
$ I = 3$	40	80	72	mixed Zg quark currents and the two ordered $J_d(2, 4, 5)$ families; current IDs 29–68
$ I = 4$	48	144	120	endpoint currents with stored orders $(2, 5, 4, 3)$ and $(2, 4, 5, 3)$; current IDs 69–116
amplitude	—	48 roots	48 outputs	24 nonzero helicities for each physical colour flow
total	117	242	210	no auxiliary current is needed for this process

An *attachment* says that one interaction contributes to one target current. It need not imply a fresh kernel call. For example, the source quark current combined with $J_g(4, 5)$ can feed both ordered three-leg targets. The two attachments share one verified `evaluation_group_id`; a sign or scalar `evaluation_factor` is applied at each target. This is why the three- and four-leg layers contain 224 attachments but only 192 distinct kernel evaluations. The resulting sharing is stronger than merely caching a list of Feynman diagrams.

9.4 Coverage slices, reduction, and reproduction

Normal public generation materializes both flows and complete helicity metadata. Runtime selection then uses the stable physical IDs. The report’s specialized one-flow workload instead uses the private developer helper `GenerationSlice`: the source-label order $(2, 5, 4, 1, 3)$ identifies the coloured traversal $(2, 5, 4, 1)$, while the singlet 3 is retained only as source-order provenance. The reference flow is inserted first, so artifact-local sector 0 can be retained during that specialized generation. This sliced graph has 65 currents, 90 interactions, and 24 roots; those internal numbers are not a substitute for the public flow ID.

name space	example	meaning
source label	4	fourth particle in the user process
all-outgoing leg	$g(q_4)$	crossed DAG input with $q_4 = p_4$
internal sector ID	0	artifact-local generation index
physical flow ID	flow:2,5,4,1	stable runtime selector
current ID	69	dense slot for one complete current identity
value slot	propagated current 69	runtime storage variant and component range
root/output	one (h, c) contribution	complex amplitude before squaring

For the built-in compatibility model, Rusticol applies

$$|\mathcal{M}|^2 = \mathcal{N} \sum_{h,c} w_{h,c} \left| \sum_{r \in G(h,c)} \mathcal{A}_{h,c}^{(r)} \right|^2, \quad \mathcal{N} = \frac{27 (4\pi\alpha_s)^2 (8\pi\alpha_{EW})}{36 \times 2}. \quad (14)$$

The denominator averages the two incoming colour/spin state spaces (6×6); the final factor $2 = 2!$ accounts for identical gluons. The physics manifest records the factors, coupling powers, coherent groups, and symmetry weights separately, so changing runtime model parameters refreshes the normalization without rebuilding evaluator code.

The installed examples provide a checkout-independent public route:

```
pyamplicol examples copy ./pyamplicol-examples
cd pyamplicol-examples
pyamplicol generate_pp_zjj_from_ufo_sm.toml
pyamplicol inspect artifacts/pp_zjj --process 'd d~ > z g g'
pyamplicol profile artifacts/pp_zjj --process 'd d~ > z g g' \
  --target-runtime 1 --batch-size 128
```

The generated artifact contains a process set; the exact expression is therefore the clearest selector. In Python, the same complete artifact exposes both the summed and resolved contracts:

```
from pyamplicol import Runtime

runtime = Runtime.load("artifacts/pp_zjj", process="d d~ > z g g")
print(runtime.physics.color_flow_ids)
total = runtime.evaluate(momenta)
resolved = runtime.evaluate_resolved(
    momenta, color_flows=["flow:2,5,4,1"]
)
```

The developer-only specialization lives in `tools/developer/generation_slice.py`; it is intentionally absent from the public TOML and CLI contracts. The generation and runtime APIs otherwise use the same process metadata and numerical reduction.

10 The $d\bar{d} \rightarrow Zgg$ DAG and Rusticol

The previous section described the physics graph. This section gives the exact mapping from that graph to compiled multi-output evaluators and then to the Rusticol runtime. The main invariant is simple: generation owns physics identity and symbolic lowering; Rusticol owns a validated numerical schedule of slots, stages, roots, and reductions.

10.1 What identifies a current?

The immutable `CurrentIndex` key is defined in `src/pyamplicol/generation/dag_types.py`. Two contributions share a current if and only if every field in

$$\Omega(J) = (t, m_I, I, H, \chi, s, F, Q, C, m_p, O, a, \pi_I) \quad (15)$$

is equal. There is no process-family or diagram label in this key.

field	example	role
<code>particle_id</code> t	1	off-shell particle carried by the current
<code>external_mask</code> m_I	26	bit mask of source subset $I = \{2, 4, 5\}$
<code>external_labels</code> I	(2, 4, 5)	sorted, unordered source-label set
<code>helicity_ancestry</code> H	source-state bit set	exact one-leg state choices feeding the current
<code>chirality</code> χ	+1 or -1	Weyl-flow branch when applicable
<code>spin_state</code> s	scalar or tuple	local spin/polarization state beyond chirality
<code>flavour_flow</code> F	model tuple	conserved flavour-line identity
<code>quantum_number_flow</code> Q	name/expression pairs	exact model-derived additive charges
<code>color_state</code> C	LC state	accuracy, internal sector, line groups, and basis keys
<code>momentum_mask</code> m_p	26	momentum slot used by kernels and propagator
<code>coupling_orders</code> O	QCD=2	accumulated perturbative order envelope
<code>auxiliary_kind</code> a	None	proof-backed contact-decomposition species, if present
<code>ordered_external_labels</code>	(2, 5, 4)	colour-sensitive traversal within the same subset
π_I		

Both the unordered labels and their order are intentional: the former fixes the momentum and subset, while the latter prevents two colour-ordered currents from being merged incorrectly. Conversely, exact equality permits reuse across every diagram and amplitude root that reaches the same state. The `_CurrentTable.add_or_get()` operation in `src/pyamplicol/generation/dag_table.py` is the sole identity-deduplication point. The resulting `current_id` is only the dense array index of that canonical node.

LC generation adds a second, complementary reuse mechanism. Internal currents are built in a shared-ordering state and their ordered labels and basis keys record which flow segments they can serve. Physical sector IDs are restored only when a colour line closes into an amplitude root. The two roots in Eq. (6) therefore share sources, electroweak subcurrents, the $J_g(4, 5)$ family, and every other exactly equal current, instead of cloning a complete recursion per flow.

Each one-leg state receives its own bit when sources are constructed; composing currents takes the bitwise OR. Thus H identifies concrete source-state choices, rather than merely the set of external labels.

Finally, `evaluation_group_id` addresses reuse that current identity alone cannot express. If one verified kernel value contributes, possibly with different scalar `evaluation_factors`, to several target currents, the symbolic stage computes it once. For the worked graph this reduces 242 target attachments to 210 kernel evaluations. These three notions must remain distinct:

$$\underbrace{\text{same current key,}}_{\text{one stored node}} \quad \underbrace{\text{same evaluation group,}}_{\text{one symbolic kernel value}} \quad \underbrace{\text{same physical flow}}_{\text{one public output axis member}}. \quad (16)$$

10.2 Topological stages and value slots

The compiler sweeps proper source subsets in increasing cardinality. An interaction whose result covers k source labels belongs to recursion stage k . No evaluator can read a current from its own or a later stage, so the schedule is acyclic by construction. For this example,

$$\boxed{11 \text{ sources}} \longrightarrow \boxed{|I| = 2 : 18} \longrightarrow \boxed{|I| = 3 : 40} \longrightarrow \boxed{|I| = 4 : 48} \longrightarrow \boxed{48 \text{ amplitude roots}}. \quad (17)$$

For one concrete source-state path, the same schedule can be read as

$$E_2 : \quad C_{21} = \mathcal{P}_{21} K_{10}(C_2, C_9), \quad C_{25} = \mathcal{P}_{25} K_{14}(C_7, C_9), \quad (18)$$

$$E_3 : \quad C_{53} = \mathcal{P}_{53} [K_{66}(C_{21}, C_7) + K_{82}(C_2, C_{25}) + \dots], \quad (19)$$

$$E_4 : \quad C_{69} = \mathcal{P}_{69} [K_{90}(C_{53}, C_4) + \dots], \quad (20)$$

$$E_{\text{root}} : \quad \mathcal{A}_0 = \mathcal{C}_0(C_{69}, C_1). \quad (21)$$

Here C_2 is the label-2 quark source with helicity -1 , C_7 and C_9 are the negative-helicity gluon sources on labels 4 and 5, C_4 is the negative-helicity Z source, and C_1 is the endpoint antiquark source. The omitted terms are other admissible partitions contributing coherently to the same target. The subscripts on K are evaluation-group IDs, not current or interaction IDs.

```
sources: C2=d(2,-), C7=g(4,-), C9=g(5,-), C4=Z(3,-), C1=dbar(1,+)
E2: I10: C2 + C9 -> C21 [G10] I14: C7 + C9 -> C25 [G14]
E3: I66: C21 + C7 -> C53 [G66] I83: C2 + C25 -> C53 [G82]
      I82: C2 + C25 -> C61 [G82] (one kernel value, two targets)
E4: I98: C53 + C4 -> C69 [G90]
root: R0: contract(C69,C1) -> A0
```

This adjacency-list excerpt is diagnostic but exact for the current graph. It also exposes why the namespaces cannot be conflated: C_{53} is a canonical current node, I_{83} is one attachment to it, G_{82} identifies a reusable symbolic evaluation, and R_0 is an amplitude output. Value-slot IDs are a further runtime namespace assigning propagated or unpropagated component storage to currents. Internal LC sector IDs annotate root construction; public IDs such as `flow:2,5,4,1` label the resolved physics axis.

The stage records are assembled by `build_runtime_schema()` in `src/pyamplicol/generation/runtime_schema.py`. They list input/output current IDs, value-slot IDs, momentum slots, interaction IDs, and both attachment and unique-evaluation counts.

A current and its stored value are not identical notions. Depending on its consumers, one current can have a `source`, `unpropagated`, or `propagated` value slot. Each slot owns a contiguous component range; Weyl currents have two complex components and vector currents have four in the worked model. Keeping variants explicit prevents the runtime from silently applying a propagator twice and lets the amplitude closure request precisely the form it needs.

The complete schema for this graph has 256 complex value components, 13 four-vector momentum slots (52 real components), and 16 real model-parameter slots. A conceptual flattened state row therefore has

$$x = [x_{\text{value}} \mid x_{\text{momentum}} \mid x_{\text{model}}], \quad |x| = 256 + 52 + 16 = 324. \quad (22)$$

No compiled stage accepts all 324 entries. Stage-local layout is mandatory: for stage k , a map L_k selects only the current components, momentum components, and model parameters actually used there,

$$x_k = L_k x, \quad y_k = E_k(x_k), \quad x[S_k] \leftarrow y_k. \quad (23)$$

The construction is in `src/pyamplicol/generation/stage_parameters.py`; Rusticol stores the same local-to-global indices in the execution manifest. Contiguous index runs are precomputed, so most packing and output assignment becomes slice copies rather than scalar lookup loops.

10.3 What Symbolica and SymJIT compile

`build_generic_stage_compiler_blueprint()` in `src/pyamplicol/generation/stage_planning.py` creates one multi-output symbolic blueprint per recursion stage and one for the amplitude roots. For a target current j , its output expression is

$$y_j = \mathcal{P}_j \left[\sum_{e: \text{target}(e)=j} f_e K_{g(e)}(x_k) \right], \quad (24)$$

where $K_{g(e)}$ is cached by evaluation-group ID, f_e is the attachment factor, and $\mathcal{P}_j$ is the output propagator or identity. Component expressions, couplings, momentum dependence, and propagators come from the selected model; the stage compiler does not know that this example contains a Z or a gluon.

Symbolica optimizes the whole multi-output expression of one stage: common subexpressions, Horner schemes, and common-pair elimination can therefore cross current-component boundaries inside that stage. With the default JIT backend, SymJIT O3 then compiles a batched complex-f64 application. The direct application capability is recorded as `symjit.application.complex-f64.v1`; Rusticol can load and execute this f64 payload without importing Symbolica. Retained Symbolica evaluator state is used for higher-precision Python requests.

Output chunking is a compilation boundary, not a physics boundary. If a stage has more than the configured 128 scalar outputs, it is split into multi-output chunks sharing the same local input vector. A fanout-aware ordering in `src/pyamplicol/generation/stage_planning.py` places outputs with common evaluation groups near one another, reducing duplicated work when a group would otherwise cross chunk boundaries. At runtime the local input is packed once per stage and batch, each compiled chunk consumes that same packed buffer, and the concatenated outputs are scattered to their recorded slots. Dependencies are still cut only between recursion stages.

Generation streams these stages. Once a blueprint has been compiled and written, its Symbolica expressions can be released before the next stage is built. This bounds peak expression memory without changing the DAG or the runtime schedule. The relevant writer path is `src/pyamplicol/generation/stage_artifacts.py`, while backend settings and payload construction live in `src/pyamplicol/evaluators/symbolica_compile.py`.

10.4 Artifact records and Rusticol execution

The root `artifact.json` is a process-artifact schema-v3 manifest. It binds global model/configuration data and payload hashes to process-local records. The files relevant to one concrete process are:

record	runtime meaning
<code>processes/evaluators.json</code>	evaluator capabilities and payload locations
<code>processes/<id>/execution.json</code>	value/momentum/model layout, sources, ordered stages, chunks, roots, and reduction schedule
<code>processes/<id>/physics.json</code>	external particles, physical helicity and flow IDs, coverage, symmetry expansion, selectors, and normalization
<code>evaluator applications/libraries</code>	target-specific compiled numerical kernels
<code>API/</code>	validation point and Python, Rust, C++, and Fortran examples

The complete `CurrentIndex` is a generation contract. The execution projection serializes only the fields needed to validate and run its slots; ordered-label proofs and other generation-only evidence need not burden the hot runtime. By contrast, physical helicity and flow metadata remains explicit because it defines the public resolved axes and selector semantics.

Rusticol’s f64 path is implemented in `rust/crates/rusticol-core/src/engine/evaluation.rs`. For a batch of points, one execution is:

1. reuse and clear the state scratch buffer;
2. fill source wavefunctions in their source value slots;
3. form all recorded all-outgoing momentum sums and copy current model parameter values into the global state;
4. for every recursion stage, pack x_k , call all evaluator chunks, and assign y_k to the global value slots;
5. pack and call the amplitude-root evaluator;
6. sum roots coherently, contract colour, expand exact physical symmetry representatives when resolved output is requested, and apply normalization.

The stage and amplitude implementations reuse their parameter and output scratch buffers across calls. The global model-parameter vector is mutable at runtime; an atomic update validates all supplied values, refreshes derived parameters and masses, and recomputes the factor in Eq. (14). No mass, width, or coupling is baked into the orchestration schedule merely because it is constant in one run card.

For LC resolved evaluation, the public tensor has shape

$$R_{n,h,c} = |\mathcal{M}_{h,c}(p_n)|^2, \quad \text{shape}(R) = (N_{\text{point}}, N_{\text{helicity}}, N_{\text{flow}}). \quad (25)$$

For this process the final two dimensions have lengths 48 and 2. Structural zeros occupy their declared entries. `evaluate_resolved()` can select stable helicity/flow IDs and materializes these components; its `total()` is their explicit sum. `evaluate()` follows the optimized summed path directly and need not allocate the resolved tensor. Neither interface exposes internal current, slot, or sector IDs.

10.5 The two runtime timing boundaries

The profiler deliberately distinguishes orchestration from compiled evaluator calls. Its reported f64 wall sample starts in the native Rusticol method after the caller-language buffer conversion and repeatedly executes the ordinary core path through final reduction. It therefore includes Rust-side input batch preparation, source and momentum construction, stage-local packing, chunk calls, output assignment, root packing/call, and reduction, but excludes Python/NumPy, C++, Fortran, or Rust caller-side packing.

The evaluator sum is narrower:

$$t_{\text{eval}} = \sum_k t(E_k) + t(E_{\text{root}}), \quad t_{\text{core}} = t_{\text{eval}} + t_{\text{source}} + t_{\text{momentum/model}} + t_{\text{pack}} + t_{\text{assign}} + t_{\text{reduce}} + t_{\text{batch}}. \quad (26)$$

provenance marker	quantity
<code>runtime_core_repeated_wall_time</code>	native repeated total path, normalized per point; common metric for all language APIs
<code>runtime_profile_core_evaluator_call_time</code>	sum of recursion-stage and amplitude-root evaluator call times only
native component profile	source fill, momentum/model setup, per-stage input pack, evaluator call and output assignment, root pack/call, and reduction

These boundaries make a wall/evaluator difference interpretable as work inside Rusticol rather than Python probing overhead. They also make Python, Rust, C++, and Fortran measurements comparable when they call the same f64 artifact. The CLI calibration, uncertainty estimate, progress display, and interruptible sampling policy are implemented in `src/pyamplicol/benchmarking.py`; the native timers themselves are in `rust/crates/rusticol-core/src/engine/native_runtime.rs` and `rust/crates/rusticol-core/src/engine/evaluation.rs`.

11 Arbitrary UFO and Serialized-JSON Models

The external-model path is a compiler, not a permissive UFO interpreter. It accepts a trusted UFO Python module or its serialized JSON representation, normalizes the model to a typed, model-independent intermediate representation, and only then exposes interactions to the current-recursion generator. A structure is executable when the compiler can construct its component expressions and prove every transformation used to lower or reuse them. Failure to prove a useful optimization merely disables that optimization; failure to prove a transformation required for correctness rejects the model with a structured diagnostic.

This section describes the implemented contract. The handwritten Standard Model in `src/pyamplicol/models/builtin/` is deliberately outside that contract: `BuiltinSMModel` may use SM-specific formulae, while `CompiledUFOModel` must derive its decisions from particle metadata, normalized tensors, propagators, couplings, and exact algebra. The boundary is audited by `tests/unit/test_model_boundary_audit.py`; in particular, SM PDG classifiers and built-in implementation symbols may not leak into the external-model modules.

11.1 One route from UFO and JSON to an executable model

The public entry point `compile_model_source()` in `src/pyamplicol/models/loading.py` recognizes a UFO directory, a JSON model, an already compiled model, or the special built-in model. For an external source it lazily invokes `ufo-model-loader`. Both source forms therefore enter the same route:

$$\begin{aligned} \{\text{UFO directory, serialized JSON}\} &\xrightarrow{\text{load_model}} \text{ufo-model-loader.Model} \\ &\xrightarrow{\text{compact serialization and compilation}} \text{CompiledModelIR}. \end{aligned} \quad (27)$$

The UFO directory is executable Python and must therefore be trusted. The JSON form is data-only and is the portable form to archive or exchange.

The loader is first called with wrapped Lorentz indices disabled. This lets `_classify_external_propagators()` compare a model-supplied propagator against the loader’s exact default Feynman-gauge candidates. Indices are then wrapped and the model is serialized with `JSONLook.COMPACT`. Consequently, UFO and JSON inputs do not maintain separate physics implementations after loading.

Every source symbol is put in a model-local namespace by `src/pyamplicol/_internal/physics/symbols.py`. A model called `sm` uses heads such as `model_sm::GC_1`, while compiler-owned currents, momenta, dummy indices, and derived couplings use the `pyamplicol::` namespace. The registry rewrites the loader’s temporary `UFO::` heads; a compiled `CompiledModelIR` is rejected if an executable expression still contains a process-global `UFO::` symbol. Two loaded models can thus coexist without silently identifying equally named parameters or tensors.

The canonical IR in `src/pyamplicol/models/contracts.py` contains typed records for particles, parameters, coupling orders, propagators, vertex terms, oriented trivalent kernels, direct and closure contractions, tensor and current orderings, and Goldstone partners. It is not a bag of strings. On reload it checks the model-schema version, compiler version and fingerprint, proof identities, tensor-ordering links, and required fields before any DAG is built. The compiler fingerprint includes the `pyAmpliCol`, `Symbolica`, and `ufo-model-loader` versions and the relevant serialization and lowering policies; a stale compiled model is regenerated rather than guessed compatible.

compiler stage	resulting contract	implementation
source loading	one compact loader-model payload	<code>models/loading.py: _load_external_model()</code> and <code>preflight_model()</code>
parameter resolution	namespaced external and derived expressions	<code>models/compiler_records.py</code> and <code>ModelSymbolRegistry</code>
tensor normalization	canonical colour and Lorentz expressions	<code>models/tensors.py</code> , using <code>idenso</code> , <code>Spenso</code> , and <code>Symbolica</code>
kernel construction	oriented trivalent contracts and contact auxiliaries	<code>models/compiler_entry.py</code> , <code>models/compiler_kernels.py</code> , <code>models/compiler_contacts.py</code>
ordering construction	explicit source axes, Weyl projections, and current embeddings	<code>models/compiler_tensor_ordering.py</code>
model-owned closure	propagators, direct/closure contractions, and Goldstone policy	<code>models/compiler_contractions.py</code> , <code>models/compiler_gauge.py</code>
optimization proofs	kernel equivalence, colour-flow auxiliaries, and exact model symmetries	<code>models/compiler_color_flow.py</code> and <code>models/external_symmetries.py</code>

11.2 From a UFO vertex to concrete current kernels

A UFO vertex is a matrix over colour and Lorentz structures. The compiler expands it into separate terms

$$\mathcal{V}(1, \dots, n) = \sum_{a,b} g_{ab} C_a(c_1, \dots, c_n) L_b(s_1, p_1; \dots; s_n, p_n), \quad (28)$$

and stores the original sources together with their normalized forms and the coupling-order vector of g_{ab} . The normalization in `src/pyamplicol/models/tensors.py` converts wrapped UFO heads to typed Spenso expressions. It applies idenso's `simplify_metrics`, `simplify_gamma`, and `simplify_color`; Spenso then materializes dense tensor components in a declared Cartesian order. Residual UFO tensor heads, indeterminate expressions, duplicate coordinates, or missing dense coordinates are errors.

The currently normalized Lorentz basis includes scalars, momenta, metrics, Dirac identity and gamma matrices, γ_5 , chiral projectors, slashed momenta, and sigma tensors. The supported colour representations are the SU(3) singlet, fundamental, antifundamental, and adjoint. Exact trilinear projection recognizes four executable colour families:

$$1, \quad \delta_i^{\bar{j}}, \quad (T^a)_i^{\bar{j}}, \quad f^{abc}. \quad (29)$$

For a proposed basis tensor B_σ , projection succeeds only when the normalized tensor obeys the componentwise identity

$$C_a(\mathbf{c}) = \kappa_a B_\sigma(\mathbf{c}) \quad \text{for every materialized colour component } \mathbf{c}, \quad (30)$$

with an exact scalar κ_a . Numerical projection exists as a diagnostic but does not certify an executable kernel. The symmetric adjoint invariant d^{abc} , for example, can be recognized diagnostically but is not in the current executable trilinear colour basis.

For a trilinear term, `_compile_oriented_kernels()` chooses each distinct result species and both nonidentical input orders. Spenso contracts the normalized tensor with formal left and right currents and physical momenta. The resulting component contract has the schematic form

$$K_{r \leftarrow ij}^{A_r}(J_i, J_j) = g_{ab} \kappa_a \sum_{A_i, A_j} L_b^{A_r A_i A_j}(p_i, p_j, p_r) J_i^{A_i} J_j^{A_j}, \quad p_r = -(p_i + p_j), \quad (31)$$

followed by the model-owned propagator when the result is an internal current. The coupling, perturbative orders, runtime parameters, source-leg map, colour normalization, and input/output ordering identifiers remain attached to the kernel. Weyl projection is therefore an explicit ordering/embedding contract, not an inference made later from a particle name.

The compiler recognizes a compact Yang–Mills three-vector form only after its components have been proved equal, up to one exact constant, to the dense Spenso result. This is an implementation optimization, not an alternative physics formula.

11.3 What a proof or certificate asserts

The word *proof* is used narrowly. It is a reproducible algebraic witness that an optimized or decomposed executable contract reconstructs the normalized source contract. It is not a statement that a model is physically consistent, unitary, or gauge complete. Three assurance mechanisms are currently distinct:

1. A **serialized decomposition proof** is attached to a four-point `CompiledVertexTerm`. Its algorithm name and version, source identity, split topology, index map, component basis, signs, and multiplicities are revalidated when the compiled model is loaded.
2. **Kernel-equivalence metadata** is written on oriented kernels. Fully lowered component expressions are compared under input exchange and an exact factor $\eta = \pm 1$. Equal kernels receive one `evaluation_class`, a representative input order, and a factor, so the DAG may evaluate the representative once and fan its value out.
3. A **model-level symmetry certificate** is re-derived by `derive_external_symmetry_certificates()` when a `CompiledUFOModel` is constructed. It binds a theorem to digests of the compiled particle, propagator, ordering, vertex, and executable-kernel contracts. It is not accepted as a free-form claim from the UFO source.

For two kernel expressions K and K' , reusable evaluation requires an exact relation

$$K'(x, y; p_x, p_y) = \eta K(y, x; p_y, p_x), \quad \eta \in \{-1, +1\}, \quad (32)$$

after resolving coupling aliases and canonicalizing all component expressions. For a global theorem, the digest is

$$D = \text{SHA256}(\text{canonicalJSON}\{\text{source, propagator, orderings, terms, kernels}\}). \quad (33)$$

The selectors and 64-character digests must be complete and mutually consistent. Renaming a contracted tensor dummy must not change D , but a different contraction graph must. The internal helper

`_canonical_tensor_product_expressions()`

in `src/pyamplicol/models/external_symmetries.py` calls Symbolica's `Expression.canonize_tensors()` on the contracted Lorentz, colour, and spinor indices. Its reserved replacement head is

`pyamplicol::canonical_tensor_index.`

Tests explicitly check alpha-renaming invariance and contraction-graph sensitivity.

The important fail-closed distinction is:

- if a reflection or parity certificate is absent, the generator keeps both configurations and remains correct, but performs less reuse;
- if a four-point contact cannot be lowered without changing its tensor, colour, coupling, or multiplicity, generation stops rather than silently dropping or approximating that interaction.

proof class	checked identity	optimization or lowering enabled
trilinear colour projection	Eq. (30) in the supported SU(3) basis	colour-flow interpretation of an oriented kernel
kernel evaluation equivalence	Eq. (32), including coupling and ordering contracts	one evaluator value reused by symmetry-related current contributions
four-point contact decomposition	exact colour topology, open-axis Lorentz basis, orientation signs, and multiplicity	two trivalent auxiliary-current stages with no artificial pole
fundamental Fierz identity	exact T^a trilinear projection and eligible massless adjoint-vector metadata	compiler-owned singlet-subtraction auxiliary for LC flow
vectorlike parity	exact fundamental-generator colour and gamma-current Lorentz tensor	global helicity-flip representative reuse
Yang–Mills cubic/quartic sector	exact f^{abc} , three-vector tensor, quartic basis, coupling relation, and closed reachable sector	pure-adjoint helicity-zero theorem and LC trace-reflection folding
local adjoint reflection	every relevant cubic kernel is antisymmetric under left/right exchange	current reflection phase reuse
Goldstone partner contract	quantum numbers, exact mass expression, uniqueness, and propagator policy	unitary-gauge absorption or explicit model-supplied Goldstone treatment
tensor/current ordering	explicit physical axes and component permutations	stable Spenso-to-current and Weyl component maps

11.4 Four-point contacts: exact auxiliary-current factorization

Every momentum-independent valence-four term receives a `CompiledContactDecompositionProof` before lowering. The serialized algorithm is `ufo-four-point-contact-decomposition`, version 2. A proof has status `proven` or `unsupported`; its identity includes the term and vertex identifiers, particle tuple, colour and Lorentz matrix indices, original sources, and normalized expressions. A stale version or an identity mismatch is rejected.

The record family is explicit:

- `CompiledContactDecompositionProof` owns the term identity and all proved result-leg splits;
- `CompiledContactDecompositionSplit` and `CompiledContactDummyIndexMapping` describe one topology and its shared colour index;
- `CompiledContactOrientationProof` fixes every input order and sign; and
- `CompiledContactUnsupportedReason` carries a stable failure code and context.

Two decomposition kinds are implemented.

Literal colour singlet. If all four legs are colour singlets and $C = 1$, the Lorentz tensor may be split directly. A literal singlet attached to any coloured leg is not treated as an implicit contraction; it fails with `literal-singlet-with-colored-legs`.

Product of two structure constants. The coloured gauge-contact path requires exactly

$$C^{a_1 a_2 a_3 a_4} = c f^{a_i a_j b} f^{b a_k a_r}, \quad (34)$$

up to exact permutations, with one and only one shared adjoint dummy b and a scalar prefactor c . For each possible result leg r , the compiler records the paired legs (i, j) , remaining leg k , dummy-index map, permutation parity, LC normalization powers, and the exact scalar coefficient. It then introduces a non-propagating auxiliary current

$$A_{\rho\sigma}^b[i, j] = f^{b a_i a_j} \sum_{A_i, A_j} P_{\rho\sigma; A_i A_j} J_i^{A_i} J_j^{A_j}, \quad (35)$$

$$J_r^{A_r}|_{V_4} = \frac{c s_{ijk r}}{m_r} f^{a_r b a_k} \sum_{\rho, \sigma, A_k} F^{A_r}_{\rho\sigma A_k} A_{\rho\sigma}^b J_k^{A_k}. \quad (36)$$

Here $s_{ijk r} = \pm 1$ is the recorded orientation parity and m_r is the number of indistinguishable assignments of the remaining input species. The original coupling and complete perturbative-order vector occur only in Eq. (36); Eq. (35) has unit coupling and no propagator. The factorization therefore introduces neither an unphysical pole nor a duplicated coupling/order weight.

The Lorentz split is not inferred from the name “four gluon.” Spenso first contracts the chosen input pair while leaving the result and remaining-leg axes open. `_compress_contact_components()` stores a basis of distinct components and reconstructs every omitted component by an exact zero or ± 1 map. The final kernel contracts that basis and must reproduce the source tensor component by component. This same mechanism covers an SU(3) four-vector term such as the four-gluon vertex and colour-singlet electroweak-like four-vector contacts built from metric tensors. What differs is the colour proof: `two-structure-constants` for the former and `literal-color-singlet` for the latter.

The compiler deduplicates identical contact partials and fuses compatible finals after proof validation. Unsupported diagnostics distinguish an incorrect topology from a missing feature: wrong factor count, malformed structure-constant indices, non-scalar colour prefactor, non-unique shared dummy, absent result leg, and a nontrivalent structure-constant topology all have separate reason codes in `src/pyamplicol/models/compiler_contacts.py`.

11.5 Arbitrary-valence colour-singlet contact trees

Colour-singlet contacts with valence $n \geq 4$ are lowered by `_compile_color_singlet_contact_trees()` in `src/pyamplicol/models/compiler_contact_trees.py`. The already proved momentum-independent four-point case uses the preceding path; other singlet contacts are split recursively at the midpoint into a balanced binary tree. For a set of input leaves S , the generic recursion is

$$A_S = \Phi_S(A_{S_L}, A_{S_R}), \quad S = S_L \dot{\cup} S_R, \quad J_r|_{V_n} = \frac{g}{m_r} \Psi_r(A_{S_L}, A_{S_R}), \quad (37)$$

where the internal A_S are non-propagating, g and its coupling orders appear only in the final map Ψ_r , and m_r removes identical-input assignment multiplicity. A balanced tree has $O(n)$ auxiliary operations and $O(\log n)$ depth.

For an all-scalar contact, each auxiliary carries one component and $\Phi_S(x, y) = xy$. This is the path exercised by the arbitrary-valence scalar-contact model. For non-scalar or momentum-dependent singlet contacts, an auxiliary

carries the current components and four-momentum of every physical leaf below it. At the final node, the original normalized n -point Lorentz expression is reconstructed with

$$p_r = - \sum_{i \neq r} p_i. \quad (38)$$

On Linux, very large dense contractions are evaluated in bounded stages when the direct input work exceeds 1024 components; each staged group is limited to 256 input components and its intermediate tensor is materialized before the next contraction. This changes only contraction order. The direct function `eager_color_singlet_vertex_term_components()` remains the component oracle used to test the tree result.

This is deliberately not advertised as an arbitrary coloured n -point algorithm. It is generic in valence, spin, momentum dependence, and Lorentz expression only after the colour tensor has been proved to be a singlet that requires no unresolved coloured internal channel.

11.6 Colour-flow and model-level symmetry proofs

The LC compiler derives its singlet-subtraction channel from the fundamental Fierz identity

$$(T^a)_i^{\bar{j}} (T^a)_k^{\bar{l}} = T_R \left(\delta_i^{\bar{l}} \delta_k^{\bar{j}} - \frac{1}{N_c} \delta_i^{\bar{j}} \delta_k^{\bar{l}} \right). \quad (39)$$

The compiler function

`synthesize_fundamental_fierz_auxiliaries()`

in `src/pyamplicol/models/compiler_color_flow.py` is enabled only for an exactly projected T^a trilinear with fundamental fermion lines and a self-conjugate, massless adjoint vector with a standard propagator. It creates one compiler-owned colour-singlet vector auxiliary and reuses the proved Lorentz and coupling kernel. The $1/N_c$ closure coefficient is attached in one canonical orientation. No gluon name or PDG code participates.

Four further exact recognizers in `src/pyamplicol/models/external_symmetries.py` recover the reusable structure normally associated with a gauge theory:

- A **vectorlike parity kernel** has one massless adjoint vector, a fundamental/antifundamental fermion pair, exact T^a colour, and a Lorentz tensor that is a constant multiple of the gamma-current reference. Chiral-projector deformations do not pass this proof.
- A **Yang–Mills cubic kernel** has one self-conjugate massless adjoint-vector species, a standard propagator, exact f^{abc} , and the three-vector Lorentz tensor up to an exact constant. The effective cubic coupling is retained symbolically.
- A **Yang–Mills quartic group** must uniquely span the three $ff \times gg$ basis elements of the four-vector vertex, and each coefficient must obey the exact relation to the square of the cubic coupling. Recognizing three individually familiar-looking terms is insufficient if the coupling relation fails.
- A **local adjoint-current reflection** requires every relevant cubic kernel to satisfy Eq. (32) with $\eta = -1$, exact f^{abc} , and a standard propagator.

A global Yang–Mills theorem is emitted only if the reachable trivalent compiled sector is closed: starting from the candidate adjoint source, every kernel reachable from two already reachable currents must itself be certified Yang–Mills. Contact auxiliaries are included in this traversal. One extra reachable non-Yang–Mills interaction therefore disables the global theorem rather than being ignored. A complete certificate enables the pure-adjoint helicity-zero theorem, global helicity-flip representatives, LC trace reflection, and shared single-trace basis. The ordinary DAG remains complete when the certificate is absent.

These generic proofs recover the built-in-SM current topology for the established UFO-SM validation ladder without assuming that every future UFO model must do so. The current tests compare the built-in SM and UFO-SM for electroweak, QCD, top, pure-gluon, charged-current, multi-boson, and up to three-quark-line examples. They require equal minimal coupling limits and equal current, interaction, amplitude-root, and colour-sector counts; exact kernel reuse may reduce evaluator calls but may not increase them. The tests also repeat the proof after arbitrary particle-name, PDG, and inventory-order changes. This is strong evidence for the certified classes, not a universal topology theorem for all UFO syntax.

11.7 Propagators, closures, Goldstones, and spin two

Propagators are model contracts rather than hard-coded particle identities. For standard scalar, Dirac, vector, and spin-two records, the evaluator in `src/pyamplicol/models/external_evaluation.py` implements the expected runtime-mass and runtime-width kernels, schematically

$$P_0(p) = \frac{i}{p^2 - m^2 + im\Gamma}, \quad P_{1/2}(p) = \frac{i(\gamma^\mu p_\mu + m)}{p^2 - m^2 + im\Gamma}, \quad (40)$$

$$P_1^{\mu\nu}(p) = \frac{-i}{p^2 - m^2 + im\Gamma} \left(g^{\mu\nu} - \frac{p^\mu p^\nu}{m^2} \right), \quad P_2^{\mu\nu, \rho\sigma}(p) = \frac{i \Pi^{\mu\nu, \rho\sigma}(p)}{p^2 - m^2 + im\Gamma}. \quad (41)$$

The massless-vector branch is Feynman gauge and does not use the singular $p^\mu p^\nu / m^2$ term. Weyl currents use explicit chirality projections. For spin two, Π is the de Donder massless or Fierz–Pauli massive projector represented in a 16-component ordered tensor space. Synthetic contact auxiliaries have the identity map and no propagator; the Fierz singlet-subtraction auxiliary propagates as its compiler-owned massless vector.

A custom UFO propagator is accepted when its normalized numerator and denominator can be materialized through the same Spenso ordering contract. It is applied exactly and is not silently converted to a preferred gauge. Its presence intentionally disables symmetry theorems that require the standard propagator.

`compile_contraction_records()` emits explicit scalar, Dirac, opposite Weyl-chirality, mostly-minus vector, and rank-two tensor-product-metric contractions. Rusticol consumes these records and does not infer a closure from component count. Goldstone matching in `compile_goldstone_partner_records()` similarly uses quantum-number flow and the exact resolved mass expression. A unique massive vector with a default propagator absorbs its Goldstone in the unitary-gauge path; a custom vector propagator retains the model-supplied Goldstone. Ambiguous matches fail closed.

UFO spin code 5 is therefore a real current type, not a special-case scalar: it has two ordered Lorentz axes, 16 components, tensor-product closure, and a spin-two propagator. The packaged massless scalar-gravity model validates the generic tensor route, including axis relabelling and transposed source storage. Massive spin two is accepted with an explicit experimental warning; arbitrary massive spin-two phenomenology is not yet a qualified release capability.

11.8 Implemented support boundary

The support matrix below separates an exact current capability from a known unsupported class. “Proof-gated” means that qualifying models work without an SM-specific exception; it does not mean every algebraically equivalent UFO spelling is already recognized.

area	class	status	exact boundary
source	UFO directory / JSON	supported	both lower through one loader-model payload; UFO Python must be trusted
spin	scalar, Dirac, vector	supported	standard wavefunctions, propagators, current orderings, and closures
spin	massless spin two	supported	16-component tensor path tested by the packaged scalar-gravity model
spin	massive spin two	experimental	Fierz–Pauli path exists, but broad model validation is incomplete
spin	UFO ghosts	partial	one-component metadata and ordering exist; propagating ghosts require an exactly lowerable custom propagator and are not broadly qualified
spin	Majorana/FNV, spin 3/2	unsupported	no fermion-number-violating flow or Rarita–Schwinger current contract
colour	1, 3, $\bar{3}$, 8 of SU(3)	supported	LC and current contracted NLC/full basis
colour	6, $\bar{6}$, extra groups	unsupported	no sextet or multi-group flow/contraction representation
arity 3	1, δ , T , f colour	supported	exact projection plus generic materializable Lorentz components
arity 3	d^{abc} , ϵ , K_6 , T_6	unsupported	recognized or known source heads are not sufficient to construct an executable colour contract
arity 4	colour singlet	supported	momentum-independent proof path; generic singlet-tree path for other materializable tensors

area	class	status	exact boundary
arity 4	two- f gauge colour	proof-gated	one shared adjoint dummy, trivalent topology, scalar prefactor, exact Lorentz reconstruction
arity 4	other coloured tensors	unsupported	no certified auxiliary colour channel
arity > 4	colour-singlet contact	supported	balanced no-pole tree; original coupling/order only at final node
arity > 4	coloured contact	unsupported	no general irreducible-channel factorization proof
Lorentz	idenso/Spenso basis	supported	all source tensor heads must normalize and materialize with explicit axes
Lorentz	form factors / unknown functions	unsupported	no executable and serializable function contract
propagator	standard $0, \frac{1}{2}, 1, 2$	supported	runtime mass/width and explicit ordering/closure
propagator	custom UFO tensor	proof-gated	exact Spenso lowering; gauge conversion and affected symmetry reuse are disabled
symmetry	vectorlike/Yang–Mills/reflection	proof-gated	exact normalized identities and complete reachable-sector digest

Preflight in `preflight_model()` rejects unsupported spin or colour codes, Majorana fermions, form factors, wrong known-function arities, and unknown expression functions before expensive compilation. Tensor normalization and contact lowering impose the narrower executable boundary after preflight. Thus recognizing a UFO head such as `K6` during syntax inspection does not imply that a sextet colour current can be built.

11.9 Extending the proof-gated subset

A contributor should extend contracts rather than add a process-family or PDG exception. The safe sequence is:

1. add a typed source representation and normalization rule in `models/tensors.py` or the propagator compiler;
2. define the component ordering, current embedding, direct contraction, and colour-flow meaning in `models/contracts.py` and `models/compiler_tensor_ordering.py`;
3. materialize a direct component oracle before introducing a factored or reused evaluator;
4. encode the exact identity, orientation signs, coupling/order placement, and stale-proof identity in a proof record or derived certificate;
5. test source-index renaming, particle/PDG renaming, serialization round trip, forged or stale metadata, resolved-sum equality, and comparison with the direct or higher-precision oracle.

remaining class	effort	main missing contract
additional scalar UFO functions	low–medium	define a serializable Symbolica lowering and derivatives/branch semantics where relevant
adjoint d^{abc} trilinears	medium	add a physical colour basis, projection, contraction, and LC/NLC/full interpretation
epsilon and sextet tensors	medium–high	add sextet current identities, conjugation, flow basis, and sparse colour metric before tensor lowering
other coloured four-point contacts	medium–high	prove a complete set of auxiliary irreducible channels and exact reconstruction, including repeated representations and multiplicities
coloured valence above four	high	general representation-channel tree with associativity/basis-change proofs and controlled intermediate growth
Majorana/FNV fermions	high	fermion-flow orientation, charge conjugation, signs, wavefunctions, propagators, and closure contracts
spin $3/2$	high	ordered Rarita–Schwinger wavefunctions, projectors, contractions, and tensor-component qualification
multiple colour groups	high	group-qualified representations, independent invariant tensors, product flow bases, and contracted colour metrics
general massive spin two	medium–high	broader model and gauge/constraint validation around the existing 16-component propagator contract

The principal regression suites are:

<code>tests/unit/test_contact_color_proofs.py</code>	<code>tests/unit/test_contact_tensor_execution.py</code>
<code>tests/unit/test_external_symmetry_certificates.py</code>	<code>tests/unit/test_external_model_independence.py</code>
<code>tests/unit/test_external_goldstone_contract.py</code>	<code>tests/unit/test_external_spin2_contraction.py</code>
<code>tests/unit/test_compiler_tensor_order.py</code>	

They test not only successful examples but also deliberately deformed tensors, couplings, propagators, index orderings, and stale proof records. That negative coverage is what makes the advertised subset useful: an unfamiliar model either satisfies the same explicit algebraic contract or receives a diagnostic at the boundary where a new proof is required.